

CREATE STATISTICS

What is it for?

Tomas Vondra <tomas.vondra@2ndquadrant.com>



Agenda

- Quick intro into planning and estimates.
- Estimates with correlated columns.
- CREATE STATISTICS to the rescue!
 - functional dependencies
 - ndistinct
 - MCV lists
- Future improvements.



ZIP_CODES

```
CREATE TABLE zip_codes (  
    postal_code      INT,  
    place_name       VARCHAR(180),  
    state_name       VARCHAR(100),  
    province_name    VARCHAR(100),  
    latitude         REAL,  
    longitude        REAL  
);
```

```
cat create-table.sql | psql test
```

```
cat zip-codes-ukraine.csv | psql test -c "copy zip_codes from stdin"
```

```
-- http://download.geonames.org/export/zip/
```



EXPLAIN

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Львів';
```

QUERY PLAN

```
-----
Seq Scan on zip_codes  (cost=0.00..45864.80 rows=3154 width=61)
    (actual rows=2944 loops=1)
    Filter: ((place_name)::text = 'Львів'::text)
    Rows Removed by Filter: 1889600
Planning Time: 0.125 ms
Execution Time: 166.881 ms
(5 rows)
```



reltuples , relpages

```
SELECT reltuples, relpages FROM pg_class  
WHERE relname = 'zip_codes';
```

reltuples	relpages
1.892544e+06	22208



```
SELECT * FROM pg_stats
WHERE tablename = 'zip_codes'
AND attname = 'place_name';
```

```
-----+-----
schemaname      | public
tablename       | zip_codes
attname         | place_name
...            | ...
most_common_vals | {Київ, Харків, Дніпро, ..., Львів, ...}
most_common_freqs | {0.006367, 0.003067, 0.0029, ..., 0.001667, ...}
...            | ...
```



```
EXPLAIN (ANALYZE, TIMING off)
```

```
SELECT * FROM zip_codes WHERE place_name = 'Львів';
```

QUERY PLAN

```
-----  
Seq Scan on zip_codes  (cost=0.00..45864.80 rows=3154 width=61)
```

```
      (actual rows=2944 loops=1)
```

```
    Filter: ((place_name)::text = 'Львів'::text)
```

```
    Rows Removed by Filter: 1889600
```

```
reltuples           | 1.892544e+06
```

```
most_common_vals    | {..., Львів, ...}
```

```
most_common_freqs   | {..., 0.001667, ...}
```

$$0.001667 * 1.892544e+06 = 3154$$



```
EXPLAIN (ANALYZE, TIMING off)
```

```
SELECT * FROM zip_codes WHERE province_name = 'Lvivska';
```

QUERY PLAN

```
-----  
Seq Scan on zip_codes  (cost=0.00..45864.80 rows=2870 width=61)
```

```
      (actual rows=4288 loops=1)
```

```
    Filter: ((province_name)::text = 'Lvivska'::text)
```

```
    Rows Removed by Filter: 1888256
```

```
reltuples           | 1.892544e+06
```

```
most_common_vals    | {..., Lvivska, ...}
```

```
most_common_freqs   | {..., , 0.0015165, ...}
```

$$1.892544e+06 * 0.0015165 = 2870$$



Underestimate

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Львів'
        AND province_name = 'Lvivska';
```

QUERY PLAN

```
Seq Scan on zip_codes  (cost=0.00..50596.16 rows=5 width=61)
    (actual rows=2944 loops=1)
  Filter: (((place_name)::text = 'Львів'::text)
    AND ((province_name)::text = 'Lvivska'::text))
Rows Removed by Filter: 1889600
```



$$P(A \& B) = P(A) * P(B)$$



```
SELECT * FROM zip_codes
    WHERE place_name = 'Львів'
    AND province_name = 'Lvivska';
```

```
P(place_name = 'Львів' & province_name = 'Lvivska')
= P(place_name = 'Львів') * P(province_name = 'Lvivska')
= 0.001667 * 0.0015165
= 0.0000025280055
```

$$0.0000025280055 * 1.892544e+06 = 4.78$$



Underestimate

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Львів'
        AND province_name = 'Lvivska';
```

QUERY PLAN

```
Seq Scan on zip_codes  (cost=0.00..50596.16 rows=5 width=61)
    (actual rows=2944 loops=1)
  Filter: (((place_name)::text = 'Львів'::text)
    AND ((province_name)::text = 'Lvivska'::text))
Rows Removed by Filter: 1889600
```



Overestimate

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Київ'
        AND province_name != 'Kyiv';
```

QUERY PLAN

```
Seq Scan on zip_codes  (cost=0.00..50596.16 rows=11972 width=61)
    (actual rows=0 loops=1)
   Filter: (((province_name)::text <> 'Kyiv'::text)
            AND ((place_name)::text = 'Київ'::text))
  Rows Removed by Filter: 1892544
```



Correlated columns

- Attribute Value Independence Assumption (AVIA)
 - may result in wildly inaccurate estimates
 - both underestimates and overestimates
- consequences
 - poor scan choices (Seq Scan vs. Index Scan)
 - poor join choices (Nested Loop)



Poor scan choices

Index Scan using orders_city_idx on orders

(cost=0.28..185.10 **rows=90** width=36)

(actual **rows=12248237** loops=1)

Seq Scan using on orders

(cost=0.13..129385.10 **rows=12248237** width=36)

(actual **rows=90** loops=1)



Poor join choices

```
-> Nested Loop (... rows=90 ...) (... rows=12248237 ...)  
    -> Index Scan using orders_city_idx on orders  
        (cost=0.28..185.10 rows=90 width=36)  
        (actual rows=12248237 loops=1)  
        ...  
    -> Index Scan ... (... loops=12248237)
```




Poor join choices

```
-> Nested Loop (... rows=90 ...) (... rows=12248237 ...)
-> Nested Loop (... rows=90 ...) (... rows=12248237 ...)
  -> Nested Loop (... rows=90 ...) (... rows=12248237 ...)
    -> Index Scan using orders_city_idx on orders
        (cost=0.28..185.10 rows=90 width=36)
        (actual rows=12248237 loops=1)
        ...
    -> Index Scan ... (... loops=12248237)
  -> Index Scan ... (... loops=12248237)
-> Index Scan ... (... loops=12248237)
```



functional dependencies (WHERE)



Functional Dependencies

- value in column A determines value in column B
- trivial example: primary key determines everything
 - zip code \rightarrow {place, state, province}
 - 4176 \rightarrow {Київ, Kyiv, Kyiv}
- other dependencies:
 - place \rightarrow state
 - state \rightarrow province
 - place \rightarrow province



CREATE STATISTICS

```
CREATE STATISTICS s (dependencies)
    ON place_name, state_name, province_name FROM zip_codes;
    2           3           4
ANALYZE zip_codes;
SELECT stxdependencies FROM pg_statistic_ext WHERE stxname = 's';
```

stxdependencies

```
-----
{"2 => 3": 0.683733, "2 => 4": 0.658233,
 "3 => 2": 0.005067, "3 => 4": 0.005067,
 "4 => 2": 0.029000, "4 => 3": 0.949033,
 "2, 3 => 4": 0.862367, "2, 4 => 3": 0.999400, "3, 4 => 2": 0.029000}
```



place → province: 0.658233 = d

$$P(\text{place} = \text{'Львів'} \ \& \ \text{province} = \text{'Lvivska'}) = \\ P(\text{place} = \text{'Львів'}) * [d + (1-d) * P(\text{province} = \text{'Lvivska'})]$$
$$1.892544\text{e}+06 * 0.001667 * (0.658233 + (1.0 - 0.658233) * 0.0015165) \\ = 2078.2$$



Underestimate : fixed

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Львів'
        AND province_name = 'Lvivska';
```

QUERY PLAN

```
Seq Scan on zip_codes  (cost=0.00..50596.16 rows=2078 width=61)
    (actual rows=2944 loops=1)
   Filter: (((place_name)::text = 'Львів'::text)
            AND ((province_name)::text = 'Lvivska'::text))
  Rows Removed by Filter: 1889600
```



Overestimate #1: not fixed :-)

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Київ'
        AND province_name != 'Kyiv';
```

QUERY PLAN

```
Seq Scan on zip_codes  (cost=0.00..50596.16 rows=11972 width=61)
    (actual rows=0 loops=1)
   Filter: (((province_name)::text <> 'Kyiv'::text)
            AND ((place_name)::text = 'Київ'::text))
  Rows Removed by Filter: 1892544
```

Functional dependencies only work with equalities.



Overestimate #2: not fixed :-)

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Львів'
        AND province_name = 'Kyiv';
```

QUERY PLAN

```
Seq Scan on zip_codes  (cost=0.00..50596.16 rows=2415 width=61)
    (actual rows=0 loops=1)
   Filter: (((place_name)::text = 'Львів'::text) AND
((province_name)::text = 'Kyiv'::text))
  Rows Removed by Filter: 1892544
```

The queries need to respect the functional dependencies.



ndistinct (GROUP BY)



```
EXPLAIN (ANALYZE, TIMING off)
```

```
SELECT count(*) FROM zip_codes GROUP BY province_name;
```

QUERY PLAN

```
-----  
HashAggregate  (cost=50596.16..50601.38 rows=522 width=20)
```

```
    (actual rows=524 loops=1)
```

```
    Group Key: province_name
```

```
    -> Seq Scan on zip_codes  (cost=0.00..41133.44 rows=1892544 width=12)
```

```
        (actual rows=1892544 loops=1)
```

```
Planning Time: 0.111 ms
```

```
Execution Time: 382.878 ms
```

```
(5 rows)
```



```
SELECT attname, n_distinct
FROM pg_stats WHERE tablename = 'zip_codes';
```

attname		n_distinct
state_name		25
postal_code		25274
place_name		16788
latitude		12154
longitude		13694
province_name		522

(7 rows)



```
EXPLAIN (ANALYZE, TIMING off)
```

```
SELECT count(*) FROM zip_codes GROUP BY place_name, province_name;
```

QUERY PLAN

```
-----  
GroupAggregate  (cost=341951.09..362769.07 rows=189254 width=46)
```

```
    (actual rows=27631 loops=1)
```

```
    Group Key: place_name, province_name
```

```
    -> Sort  (cost=341951.09..346682.45 rows=1892544 width=38)
```

```
        (actual rows=1892544 loops=1)
```

```
        Sort Key: place_name, province_name
```

```
        Sort Method: external merge  Disk: 91360kB
```

```
        -> Seq Scan on zip_codes  (cost=0.00..41133.44 rows=1892544 width=38)
```

```
            (actual rows=1892544 loops=1)
```



```
EXPLAIN (ANALYZE, TIMING off)
```

```
SELECT count(*) FROM zip_codes GROUP BY place_name, province_name;
```

QUERY PLAN

```
GroupAggregate  (cost=341951.09..362769.07 rows=189254 width=46)
```

```
    (actual rows=27631 loops=1)
```

```
    Group Key: place_name, province_name
```

```
    ->  Sort  (cost=341951.09..346682.45 rows=1892544 width=38)
```

```
        (actual rows=1892544 loops=1)
```

```
        Sort Key: place_name, province_name
```

```
        Sort Method: external merge  Disk: 91360kB
```

```
        ->  Seq Scan on zip_codes  (cost=0.00..41133.44 rows=1892544 width=38)
```

```
            (actual rows=1892544 loops=1)
```



```
ndistinct(place, province)
=
ndistinct(place) * ndistinct(province)

16788 * 522 = 8763336
```



`ndistinct(place, province)`
=
`ndistinct(place) * ndistinct(province)`

`16788 * 522 = 8763336`

(capped to 10% of the table)



```
CREATE STATISTICS s (ndistinct)
  ON place_name, province_name, state_name
  FROM zip_codes;
```

```
ANALYZE zip_codes;
```

```
SELECT stxndistinct FROM pg_statistic_ext;
```

stxndistinct

```
{"2, 3": 22889, "2, 4": 26759, "3, 4": 536, "2, 3, 4": 26783}
```




```
EXPLAIN (ANALYZE, TIMING off)
```

```
SELECT count(*) FROM zip_codes GROUP BY place_name, community_name;
```

QUERY PLAN

```
HashAggregate  (cost=55327.52..55595.11 rows=26759 width=46)
```

```
    (actual rows=27631 loops=1)
```

```
    Group Key: place_name, province_name
```

```
    -> Seq Scan on zip_codes  (cost=0.00..41133.44 rows=1892544 width=38)
```

```
        (actual rows=1892544 loops=1)
```



ndistinct

- the “old behavior” was defensive
 - unreliable estimates with multiple columns
 - HashAggregate can’t spill to disk (OOM)
 - rather than crash do Sort+GroupAggregate (slow)
- ndistinct coefficients
 - make multi-column ndistinct estimates more reliable
 - reduced danger of OOM
 - large tables + GROUP BY multiple columns



Future Improvements

- additional types of statistics
 - MCV lists (PG12), histograms (??), ...
- statistics on expressions
 - currently only simple column references
 - alternative to functional indexes
- improving join estimates
 - using MCV lists
 - special multi-table statistics (syntax already supports it)



Questions?

Tomas Vondra

tomas.vondra@2ndquadrant.com

tomas@pgaddict.com



@fuzzycz