

# CREATE STATISTICS

*What is it for?*

**Tomas Vondra** <[tomas.vondra@2ndquadrant.com](mailto:tomas.vondra@2ndquadrant.com)>



# Agenda

- Quick intro into planning and estimates.
- Estimates with correlated columns.
- CREATE STATISTICS to the rescue!
  - functional dependencies
  - ndistinct
  - MCV lists
- Future improvements.



# ZIP\_CODES

```
CREATE TABLE zip_codes (
```

```
    postal_code      INT,  
    place_name       VARCHAR(180),  
    state_name       VARCHAR(100),  
    province_name    VARCHAR(100),  
    community_name    VARCHAR(100),  
    latitude         REAL,  
    longitude        REAL
```

```
);
```

```
cat create-table.sql | psql test
```

```
cat zip-codes-belgium.csv | psql test -c "copy zip_codes from stdin"
```

```
-- http://download.geonames.org/export/zip/
```



# EXPLAIN

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Antwerpen';
```

## QUERY PLAN

```
-----
Seq Scan on zip_codes  (cost=0.00..40796.40 rows=6935 width=56)
    (actual rows=7168 loops=1)
    Filter: ((place_name)::text = 'Antwerpen'::text)
    Rows Removed by Filter: 1726464
Planning Time: 0.335 ms
Execution Time: 156.506 ms
```



## reltuples , relpages

```
SELECT reltuples, relpages FROM pg_class  
WHERE relname = 'zip_codes';
```

reltuples	relpages
1.733632e+06	19126



```
SELECT * FROM pg_stats
WHERE tablename = 'zip_codes'
AND attname = 'place_name';
```

```
-----+-----
schemaname      | public
tablename       | zip_codes
attname         | place_name
...            | ...
most_common_vals | {Tournai, Bruxelles, Antwerpen, Mons, ...}
most_common_freqs | {0.005633333, 0.005333333, 0.004, 0.003433333, ...}
...            | ...
```



```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Antwerpen';
```

#### QUERY PLAN

```
-----
Seq Scan on zip_codes  (cost=0.00..40796.40 rows=6935 width=56)
      (actual rows=7168 loops=1)
```

```
    Filter: ((place_name)::text = 'Antwerpen'::text)
```

```
reltuples          | 1733632
most_common_vals    | {..., Antwerpen, ...}
most_common_freqs   | {..., 0.004, ...}
```

$$1733632 * 0.004 = 6935$$



```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE community_name = 'Antwerpen';
```

#### QUERY PLAN

```
-----
Seq Scan on zip_codes  (cost=0.00..40796.40 rows=39353 width=56)
    (actual rows=38912 loops=1)
    Filter: ((community_name)::text = 'Antwerpen'::text)
    Rows Removed by Filter: 1694720
```

```
reltuples      | 1733632
most_common_vals | {..., Antwerpen, ...}
most_common_freqs | {..., , 0.0226, ...}
```

$$1733632 * 0.0226 = 39353$$





# Underestimate

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Antwerpen'
        AND community_name = 'Antwerpen';
```

## QUERY PLAN

```
-----
Seq Scan on zip_codes  (cost=0.00..45130.48 rows=157 width=56)
    (actual rows=7168 loops=1)
    Filter: (((place_name)::text = 'Antwerpen'::text) AND
((community_name)::text = 'Antwerpen'::text))
    Rows Removed by Filter: 1726464
```



$$P(A \ \& \ B) = P(A) * P(B)$$



```
SELECT * FROM zip_codes
    WHERE place_name = 'Antwerpen'
    AND community_name = 'Antwerpen';
```

```
P(place_name = 'Antwerpen' & community_name = 'Antwerpen')
= P(place_name = 'Antwerpen') * P(community_name = 'Antwerpen')
= 0.004 * 0.0226
= 0.0000904
```

$1733632 * 0.0000904 = 156$



# Underestimate

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Antwerpen'
        AND community_name = 'Antwerpen';
```

## QUERY PLAN

```
-----
Seq Scan on zip_codes  (cost=0.00..45130.48 rows=157 width=56)
    (actual rows=7168 loops=1)
    Filter: (((place_name)::text = 'Antwerpen'::text) AND
((community_name)::text = 'Antwerpen'::text))
    Rows Removed by Filter: 1726464
```



# Overestimate

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Antwerpen'
        AND community_name != 'Antwerpen';
```

## QUERY PLAN

```
-----
Seq Scan on zip_codes  (cost=0.00..45130.48 rows=6777 width=56)
    (actual rows=0 loops=1)
    Filter: (((community_name)::text <> 'Antwerpen'::text) AND
((place_name)::text = 'Antwerpen'::text))
    Rows Removed by Filter: 1733632
```



## Correlated columns

- Attribute Value Independence Assumption (AVIA)
  - may result in wildly inaccurate estimates
  - both underestimates and overestimates
- consequences
  - poor scan choices (Seq Scan vs. Index Scan)
  - poor join choices (Nested Loop)



## Poor scan choices

Index Scan using orders\_city\_idx on orders  
(cost=0.28..185.10 **rows=90** width=36)  
(actual **rows=12248237** loops=1)

Seq Scan using on orders  
(cost=0.13..129385.10 **rows=12248237** width=36)  
(actual **rows=90** loops=1)



## Poor join choices

```
-> Nested Loop (... rows=90 ...) (... rows=12248237 ...)  
    -> Index Scan using orders_city_idx on orders  
        (cost=0.28..185.10 rows=90 width=36)  
        (actual rows=12248237 loops=1)  
        ...  
    -> Index Scan ... (... loops=12248237)
```





## Poor join choices

```
-> Nested Loop (... rows=90 ...) (... rows=12248237 ...)
  -> Nested Loop (... rows=90 ...) (... rows=12248237 ...)
    -> Nested Loop (... rows=90 ...) (... rows=12248237 ...)
      -> Index Scan using orders_city_idx on orders
          (cost=0.28..185.10 rows=90 width=36)
          (actual rows=12248237 loops=1)
          ...
      -> Index Scan ... (... loops=12248237)
    -> Index Scan ... (... loops=12248237)
  -> Index Scan ... (... loops=12248237)
-> Index Scan ... (... loops=12248237)
```



# functional dependencies (WHERE)



# Functional Dependencies

- value in column A determines value in column B
- trivial example: primary key determines everything
  - zip code  $\rightarrow$  {place, state, province, community}
  - 7506  $\rightarrow$  {Tournai, Wallonie, Hainaut, Tournai}
- other dependencies:
  - place  $\rightarrow$  community
  - community  $\rightarrow$  province
  - province  $\rightarrow$  state



# CREATE STATISTICS

```
CREATE STATISTICS s (dependencies)
  ON place_name, state_name, community_name FROM zip_codes;
      2           3           5
ANALYZE zip_codes;
SELECT stxdependencies FROM pg_statistic_ext WHERE stxname = 's';
```

stxdependencies

```
-----
{"2 => 3": 1.000000, "2 => 5": 1.000000,
 "3 => 5": 0.018015, "5 => 2": 0.000886,
 "5 => 3": 0.981985, "2, 3 => 5": 1.000000,
 "2, 5 => 3": 1.000000, "3, 5 => 2": 0.000886}
```



place → community: 1.0 = d

$$P(\text{place} = \text{'Antwerpen'} \ \& \ \text{state} = \text{'Antwerpen'}) =$$
$$P(\text{place} = \text{'Antwerpen'}) * [d + (1-d) * P(\text{state} = \text{'Antwerpen'})]$$

$$1733632 * 0.004 * (1.0 + (1.0 - 1.0) * 0.0226) = 6934.53$$



## Underestimate : fixed

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Antwerpen'
        AND community_name = 'Antwerpen';
```

### QUERY PLAN

---

```
Seq Scan on zip_codes  (cost=0.00..45130.48 rows=7570 width=56)
    (actual rows=7168 loops=1)
   Filter: (((place_name)::text = 'Antwerpen'::text) AND
            ((community_name)::text = 'Antwerpen'::text))
  Rows Removed by Filter: 1726464
```



## Overestimate #1: not fixed :-)

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Antwerpen'
        AND community_name != 'Antwerpen';
```

### QUERY PLAN

---

```
Seq Scan on zip_codes  (cost=0.00..45130.48 rows=7402 width=56)
    (actual rows=0 loops=1)
  Filter: (((community_name)::text <> 'Antwerpen'::text) AND
    ((place_name)::text = 'Antwerpen'::text))
Rows Removed by Filter: 1733632
```

**Functional dependencies only work with equalities.**



## Overestimate #2: not fixed :-)

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Antwerpen'
        AND community_name = 'Bruxelles';
```

### QUERY PLAN

```
-----
Seq Scan on zip_codes  (cost=0.00..45130.48 rows=7570 width=56)
    (actual rows=0 loops=1)
   Filter: (((place_name)::text = 'Antwerpen'::text) AND
            ((community_name)::text = 'Bruxelles'::text))
  Rows Removed by Filter: 1733632
```

The queries need to respect the functional dependencies.





# ndistinct (GROUP BY)



```
EXPLAIN (ANALYZE, TIMING off)
SELECT count(*) FROM zip_codes GROUP BY community_name;
```

### QUERY PLAN

```
-----
HashAggregate  (cost=45130.48..45130.95 rows=47 width=16)
    (actual rows=47 loops=1)
    Group Key: community_name
    -> Seq Scan on zip_codes  (cost=0.00..36462.32 rows=1733632 width=8)
        (actual rows=1733632 loops=1)

Planning Time: 0.111 ms
Execution Time: 382.878 ms
(5 rows)
```



```
SELECT attname, n_distinct  
FROM pg_stats WHERE tablename = 'zip_codes';
```

attname		n_distinct
-----+-----		
postal_code		1175
place_name		2847
state_name		4
province_name		12
<b>community_name</b>		<b>47</b>
latitude		496
longitude		596
(7 rows)		



```
EXPLAIN (ANALYZE, TIMING off)
```

```
SELECT count(*) FROM zip_codes GROUP BY place_name, community_name;
```

#### QUERY PLAN

```
-----  
GroupAggregate  (cost=287219.11..305893.52 rows=133809 width=32)
```

```
    (actual rows=2847 loops=1)
```

```
    Group Key: place_name, community_name
```

```
    -> Sort  (cost=287219.11..291553.19 rows=1733632 width=24)
```

```
        (actual rows=1733632 loops=1)
```

```
        Sort Key: place_name, community_name
```

```
        Sort Method: external merge  Disk: 59008kB
```

```
        -> Seq Scan on zip_codes  (cost=0.00..36462.32 rows=1733632 width=24)
```

```
            (actual rows=1733632 loops=1)
```

```
Planning Time: 0.334 ms
```

```
Execution Time: 2656.216 ms
```



```
EXPLAIN (ANALYZE, TIMING off)
```

```
SELECT count(*) FROM zip_codes GROUP BY place_name, community_name;
```

#### QUERY PLAN

```
-----  
GroupAggregate (cost=287219.11..305893.52 rows=133809 width=32)
```

```
    (actual rows=2847 loops=1)
```

```
    Group Key: place_name, community_name
```

```
    -> Sort (cost=287219.11..291553.19 rows=1733632 width=24)
```

```
        (actual rows=1733632 loops=1)
```

```
        Sort Key: place_name, community_name
```

```
        Sort Method: external merge  Disk: 59008kB
```

```
        -> Seq Scan on zip_codes (cost=0.00..36462.32 rows=1733632 width=24)
```

```
            (actual rows=1733632 loops=1)
```

```
Planning Time: 0.334 ms
```

```
Execution Time: 2656.216 ms
```



`ndistinct(place, community)`  
=  
`ndistinct(province) * ndistinct(community)`  
  
`2847 * 47 = 133809`



```
CREATE STATISTICS s (ndistinct)
  ON place_name, province_name, community_name
  FROM zip_codes;
```

```
ANALYZE zip_codes;
```

```
SELECT stxndistinct FROM pg_statistic_ext;
```

stxndistinct

---

```
{"2, 4": 2847, "2, 5": 2847, "4, 5": 48, "2, 4, 5": 2847}
```



```
EXPLAIN (ANALYZE, TIMING off)
SELECT count(*) FROM zip_codes GROUP BY place_name, community_name;
```

#### QUERY PLAN

```
-----
HashAggregate  (cost=49464.56..49493.03 rows=2847 width=32)
    (actual rows=2847 loops=1)
    Group Key: place_name, community_name
    -> Seq Scan on zip_codes  (cost=0.00..36462.32 rows=1733632 width=24)
        (actual rows=1733632 loops=1)
```

Planning Time: 0.197 ms

Execution Time: 518.319 ms





# ndistinct

- the “old behavior” was defensive
  - unreliable estimates with multiple columns
  - HashAggregate can’t spill to disk (OOM)
  - rather than crash do Sort+GroupAggregate (slow)
- ndistinct coefficients
  - make multi-column ndistinct estimates more reliable
  - reduced danger of OOM
  - large tables + GROUP BY multiple columns



# Future Improvements

- additional types of statistics
  - MCV lists, histograms, ...
- statistics on expressions
  - currently only simple column references
  - alternative to functional indexes
- improving join estimates
  - using MCV lists
  - special multi-table statistics (syntax already supports it)



Questions?

**Tomas Vondra**

[tomas.vondra@2ndquadrant.com](mailto:tomas.vondra@2ndquadrant.com)  
[tomas@pgaddict.com](mailto:tomas@pgaddict.com)



**@fuzzycz**