

# CREATE STATISTICS

*What is it for?*

Tomas Vondra, 2ndQuadrant  
<tomas.vondra@2ndquadrant.com> / <tomas@pgaddict.com>



## Agenda

- Quick intro into planning and estimates.
- Estimates with correlated columns.
- CREATE STATISTICS to the rescue!
  - functional dependencies
  - ndistinct
  - MCV lists
- Future improvements.



## ZIP\_CODES

```
CREATE TABLE zip_codes (  
    postal_code      INT,  
    place_name       VARCHAR(180),  
    state_name       VARCHAR(100),  
    province_name    VARCHAR(100),  
    latitude         REAL,  
    longitude        REAL  
);
```

```
cat create-table.sql | psql test
```

```
cat zip-codes-ukraine.csv | psql test -c "copy zip_codes from stdin"
```

```
-- http://download.geonames.org/export/zip/
```



## EXPLAIN

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Львів';
```

### QUERY PLAN

```
-----
Seq Scan on zip_codes  (cost=0.00..45864.80 rows=3154 width=61)
    (actual rows=2944 loops=1)
    Filter: ((place_name)::text = 'Львів'::text)
    Rows Removed by Filter: 1889600
Planning Time: 0.125 ms
Execution Time: 166.881 ms
(5 rows)
```



## reltuples , relpages

```
SELECT reltuples, relpages FROM pg_class  
WHERE relname = 'zip_codes';
```

reltuples	relpages
1.892544e+06	22208



```
SELECT * FROM pg_stats
WHERE tablename = 'zip_codes'
      AND attname = 'place_name';
```

```
-----+-----
schemaname      | public
tablename       | zip_codes
attname         | place_name
...             | ...
most_common_vals | {Київ, Харків, Дніпро, ..., Львів, ...}
most_common_freqs | {0.006367, 0.003067, 0.0029, ..., 0.001667, ...}
...             | ...
```



```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Львів';
```

## QUERY PLAN

```
-----
Seq Scan on zip_codes  (cost=0.00..45864.80 rows=3154 width=61)
    (actual rows=2944 loops=1)
    Filter: ((place_name)::text = 'Львів'::text)
    Rows Removed by Filter: 1889600
```

```
reltuples          | 1.892544e+06
most_common_vals   | {..., Львів, ...}
most_common_freqs  | {..., 0.001667, ...}
```

$$0.001667 * 1.892544e+06 = 3154$$



```
EXPLAIN (ANALYZE, TIMING off)
```

```
SELECT * FROM zip_codes WHERE province_name = 'Lvivska';
```

## QUERY PLAN

```
-----  
Seq Scan on zip_codes  (cost=0.00..45864.80 rows=2870 width=61)
```

```
      (actual rows=4288 loops=1)
```

```
    Filter: ((province_name)::text = 'Lvivska'::text)
```

```
    Rows Removed by Filter: 1888256
```

```
reltuples           | 1.892544e+06
```

```
most_common_vals    | {..., Lvivska, ...}
```

```
most_common_freqs   | {..., , 0.0015165, ...}
```

$$1.892544e+06 * 0.0015165 = 2870$$





## Underestimate

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Львів'
        AND province_name = 'Lvivska';
```

### QUERY PLAN

---

```
Seq Scan on zip_codes  (cost=0.00..50596.16 rows=5 width=61)
    (actual rows=2944 loops=1)
  Filter: (((place_name)::text = 'Львів'::text)
    AND ((province_name)::text = 'Lvivska'::text))
Rows Removed by Filter: 1889600
```



$$P(A \& B) = P(A) * P(B)$$



```
SELECT * FROM zip_codes
    WHERE place_name = 'Львів'
    AND province_name = 'Lvivska';
```

```
P(place_name = 'Львів' & province_name = 'Lvivska')
= P(place_name = 'Львів') * P(province_name = 'Lvivska')
= 0.001667 * 0.0015165
= 0.0000025280055
```

$$0.0000025280055 * 1.892544e+06 = 4.78$$



## Underestimate

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Львів'
        AND province_name = 'Lvivska';
```

### QUERY PLAN

---

```
Seq Scan on zip_codes  (cost=0.00..50596.16 rows=5 width=61)
    (actual rows=2944 loops=1)
   Filter: (((place_name)::text = 'Львів'::text)
            AND ((province_name)::text = 'Lvivska'::text))
  Rows Removed by Filter: 1889600
```



## Overestimate

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Київ'
        AND province_name != 'Kyiv';
```

### QUERY PLAN

---

```
Seq Scan on zip_codes  (cost=0.00..50596.16 rows=11972 width=61)
    (actual rows=0 loops=1)
  Filter: (((province_name)::text <> 'Kyiv'::text)
    AND ((place_name)::text = 'Київ'::text))
Rows Removed by Filter: 1892544
```



## Correlated columns

- Attribute Value Independence Assumption (AVIA)
  - may result in wildly inaccurate estimates
  - both underestimates and overestimates
- consequences
  - poor scan choices (Seq Scan vs. Index Scan)
  - poor join choices (Nested Loop)



## Poor scan choices

Index Scan using orders\_city\_idx on orders

(cost=0.28..185.10 **rows=90** width=36)

(actual **rows=12248237** loops=1)

Seq Scan using on orders

(cost=0.13..129385.10 **rows=12248237** width=36)

(actual **rows=90** loops=1)



## Poor join choices

```
-> Nested Loop (... rows=90 ...) (... rows=12248237 ...)  
    -> Index Scan using orders_city_idx on orders  
        (cost=0.28..185.10 rows=90 width=36)  
        (actual rows=12248237 loops=1)  
        ...  
    -> Index Scan ... (... loops=12248237)
```





## Poor join choices

```
-> Nested Loop (... rows=90 ...) (... rows=12248237 ...)
    -> Nested Loop (... rows=90 ...) (... rows=12248237 ...)
        -> Nested Loop (... rows=90 ...) (... rows=12248237 ...)
            -> Index Scan using orders_city_idx on orders
                (cost=0.28..185.10 rows=90 width=36)
                (actual rows=12248237 loops=1)
                ...
            -> Index Scan ... (... loops=12248237)
        -> Index Scan ... (... loops=12248237)
    -> Index Scan ... (... loops=12248237)
-> Index Scan ... (... loops=12248237)
```



## functional dependencies (WHERE)



## Functional Dependencies

- value in column A determines value in column B
- trivial example: primary key determines everything
  - zip code  $\rightarrow$  {place, state, province}
  - 4176  $\rightarrow$  {Київ, Kyiv, Kyiv}
- other dependencies:
  - place  $\rightarrow$  state
  - state  $\rightarrow$  province
  - place  $\rightarrow$  province



## CREATE STATISTICS

```
CREATE STATISTICS s (dependencies)
    ON place_name, state_name, province_name FROM zip_codes;
    2           3           4
ANALYZE zip_codes;
SELECT dependencies FROM pg_stats_ext WHERE statistics_name = 's';
```

dependencies

```
-----
{"2 => 3": 0.683733, "2 => 4": 0.658233,
 "3 => 2": 0.005067, "3 => 4": 0.005067,
 "4 => 2": 0.029000, "4 => 3": 0.949033,
 "2, 3 => 4": 0.862367, "2, 4 => 3": 0.999400, "3, 4 => 2": 0.029000}
```



place → province: 0.658233 = d

$$P(\text{place} = \text{'Львів'} \ \& \ \text{province} = \text{'Lvivska'}) = \\ P(\text{place} = \text{'Львів'}) * [d + (1-d) * P(\text{province} = \text{'Lvivska'})]$$
$$1.892544\text{e}+06 * 0.001667 * (0.658233 + (1.0 - 0.658233) * 0.0015165) \\ = 2078.2$$



## Underestimate : fixed

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Львів'
        AND province_name = 'Lvivska';
```

### QUERY PLAN

---

```
Seq Scan on zip_codes  (cost=0.00..50596.16 rows=2078 width=61)
    (actual rows=2944 loops=1)
  Filter: (((place_name)::text = 'Львів'::text)
    AND ((province_name)::text = 'Lvivska'::text))
Rows Removed by Filter: 1889600
```



## Overestimate #1: not fixed :-)

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE province_name = 'Kyiv'
        AND state_name != 'Kyiv';
```

### QUERY PLAN

---

```
Seq Scan on zip_codes  (cost=0.00..50131.16 rows=9290 width=60)
    (actual rows=0 loops=1)
  Filter: (((state_name)::text <> 'Kyiv'::text) AND
    ((province_name)::text = 'Kyiv'::text))
Rows Removed by Filter: 1892544
```

Functional dependencies only work with equalities.



## Overestimate #2: not fixed :-)

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE province_name = 'Kyiv'
        AND state_name = 'Lvivska';
```

### QUERY PLAN

```
-----
Seq Scan on zip_codes  (cost=0.00..50131.16 rows=10351 width=60)
    (actual rows=0 loops=1)
    Filter: (((province_name)::text = 'Kyiv'::text)
        AND ((state_name)::text = 'Lvivska'::text))
    Rows Removed by Filter: 1892544
```

The queries need to respect the functional dependencies.





## ndistinct (GROUP BY)



```
EXPLAIN (ANALYZE, TIMING off)
```

```
SELECT count(*) FROM zip_codes GROUP BY province_name;
```

## QUERY PLAN

```
-----  
HashAggregate  (cost=50596.16..50601.38 rows=522 width=20)  
    (actual rows=524 loops=1)
```

```
    Group Key: province_name
```

```
    -> Seq Scan on zip_codes  (cost=0.00..41133.44 rows=1892544 width=12)  
        (actual rows=1892544 loops=1)
```

```
Planning Time: 0.111 ms
```

```
Execution Time: 382.878 ms
```

```
(5 rows)
```



```
SELECT attname, n_distinct
FROM pg_stats WHERE tablename = 'zip_codes';
```

attname		n_distinct
state_name		25
postal_code		25274
place_name		16788
latitude		12154
longitude		13694
<b>province_name</b>	<b> </b>	<b>522</b>

(7 rows)



```
EXPLAIN (ANALYZE, TIMING off)
```

```
SELECT count(*) FROM zip_codes GROUP BY place_name, province_name;
```

## QUERY PLAN

```
-----  
GroupAggregate  (cost=341951.09..362769.07 rows=189254 width=46)
```

```
    (actual rows=27631 loops=1)
```

```
    Group Key: place_name, province_name
```

```
    -> Sort  (cost=341951.09..346682.45 rows=1892544 width=38)
```

```
        (actual rows=1892544 loops=1)
```

```
        Sort Key: place_name, province_name
```

```
        Sort Method: external merge  Disk: 91360kB
```

```
        -> Seq Scan on zip_codes  (cost=0.00..41133.44 rows=1892544 width=38)
```

```
            (actual rows=1892544 loops=1)
```



```
EXPLAIN (ANALYZE, TIMING off)
```

```
SELECT count(*) FROM zip_codes GROUP BY place_name, province_name;
```

## QUERY PLAN

```
-----  
GroupAggregate  (cost=341951.09..362769.07 rows=189254 width=46)
```

```
    (actual rows=27631 loops=1)
```

```
    Group Key: place_name, province_name
```

```
    ->  Sort  (cost=341951.09..346682.45 rows=1892544 width=38)
```

```
        (actual rows=1892544 loops=1)
```

```
        Sort Key: place_name, province_name
```

```
        Sort Method: external merge  Disk: 91360kB
```

```
        ->  Seq Scan on zip_codes  (cost=0.00..41133.44 rows=1892544 width=38)
```

```
            (actual rows=1892544 loops=1)
```



```
ndistinct(place, province)
=
ndistinct(place) * ndistinct(province)

16788 * 522 = 8763336
```



```
ndistinct(place, province)
=
ndistinct(place) * ndistinct(province)
```

```
16788 * 522 = 8763336 => 189254
```

(capped to 10% of the table)



```
CREATE STATISTICS s (ndistinct)
  ON place_name, province_name, state_name
  FROM zip_codes;
```

```
ANALYZE zip_codes;
```

```
SELECT n_distinct FROM pg_stats_ext WHERE statistics_name = 's'
```

n\_distinct

---

```
{"2, 3": 22666, "2, 4": 26573, "3, 4": 530, "2, 3, 4": 26585}
```





```
EXPLAIN (ANALYZE, TIMING off)
SELECT count(*) FROM zip_codes GROUP BY place_name, province_name;
```

## QUERY PLAN

```
-----
HashAggregate  (cost=54862.52..55128.25 rows=26573 width=46)
    (actual rows=27631 loops=1)
    Group Key: place_name, province_name
    -> Seq Scan on zip_codes  (cost=0.00..40668.44 rows=1892544 width=38)
        (actual rows=1892544 loops=1)
```



## ndistinct

- the “old behavior” was defensive
  - unreliable estimates with multiple columns
  - HashAggregate can’t spill to disk (OOM)
  - rather than crash do Sort+GroupAggregate (slow)
- ndistinct coefficients
  - make multi-column ndistinct estimates more reliable
  - reduced danger of OOM
  - large tables + GROUP BY multiple columns



## MCV lists (PG12)



```
--  
  
CREATE STATISTICS s (mcv) ON state_name, province_name FROM zip_codes;  
  
-- more detailed stats  
SET default_statistics_target = 10000;  
  
-- show the MCV list contents  
SELECT * FROM pg_mcv_list_items((  
    SELECT stxdmcv FROM pg_statistic_ext_data d  
        JOIN pg_statistic_ext s on (s.oid = d.stxoid)  
    WHERE stxname = 's'));
```



## Overestimate #1: fixed

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE province_name = 'Kyiv'
        AND state_name != 'Kyiv';
```

### QUERY PLAN

---

```
Seq Scan on zip_codes  (cost=0.00..50131.16 rows=1 width=60)
```

```
      (actual rows=0 loops=1)
```

```
    Filter: (((state_name)::text <> 'Kyiv'::text) AND
              ((province_name)::text = 'Kyiv'::text))
```

```
    Rows Removed by Filter: 1892544
```



## Overestimate #2

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE province_name = 'Kyiv'
                        AND state_name = 'Lvivska';
```

### QUERY PLAN

---

```
Seq Scan on zip_codes  (cost=0.00..50131.16 rows=1 width=60)
```

```
      (actual rows=0 loops=1)
```

```
    Filter: (((province_name)::text = 'Kyiv'::text) AND
              ((state_name)::text = 'Lvivska'::text))
```

```
    Rows Removed by Filter: 1892544
```



## summary



## Future Improvements

- additional types of statistics
  - MCV lists (PG12), histograms (??), ...
- statistics on expressions
  - currently only simple column references
  - alternative to functional indexes
- improving join estimates
  - using MCV lists
  - special multi-table statistics (syntax already supports it)





Questions?

**Tomas Vondra**

[tomas.vondra@2ndquadrant.com](mailto:tomas.vondra@2ndquadrant.com)

[tomas@pgaddict.com](mailto:tomas@pgaddict.com)



**@fuzzycz**