

CREATE STATISTICS

What is it for?

Tomas Vondra <tomas.vondra@2ndquadrant.com>



Agenda

- Quick intro into planning and estimates.
- Estimates with correlated columns.
- CREATE STATISTICS to the rescue!
 - functional dependencies
 - ndistinct
- Future improvements.



ZIP_CODES

```
CREATE TABLE zip_codes (  
    country_code      VARCHAR(2),  
    postal_code       VARCHAR(20),  
    place_name        VARCHAR(180),  
    state_name        VARCHAR(100),  
    state_code        VARCHAR(20),  
    ...  
);  
cat create.sql | psql test  
cat zip-codes-germany.csv | \  
    psql test -c "copy zip_codes from stdin \  
                  with (format csv, header true, \  
                        delimiter E'\t')"  
-- http://download.geonames.org/export/zip/
```



EXPLAIN

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Berlin';
```

QUERY PLAN

```
Seq Scan on zip_codes  (cost=0.00..57509.07 rows=22858 width=95)
    (actual rows=23296 loops=1)
```

```
    Filter: ((place_name)::text = 'Berlin'::text)
```

```
    Rows Removed by Filter: 2086144
```

```
Planning Time: 0.289 ms
```

```
Execution Time: 203.923 ms
```



reltuples , relpages

```
SELECT reltuples, relpages FROM pg_class  
WHERE relname = 'zip_codes';
```

reltuples		relpages
2.10993e+06		31135

(1 row)



```
SELECT * FROM pg_stats
WHERE tablename = 'zip_codes'
AND attname = 'place_name';
```

```
-----+-----
schemaname      | public
tablename       | zip_codes
attname         | place_name
...             | ...
most_common_vals | {Berlin, Hamburg, München, Köln, ...}
most_common_freqs | {0.0110, 0.0061, 0.0045, 0.0027, ...}
...             | ...
```



```
EXPLAIN (ANALYZE, TIMING off)
```

```
SELECT * FROM zip_codes WHERE place_name = 'Berlin';
```

QUERY PLAN

```
-----  
Seq Scan on zip_codes  (cost=0.00..57509.07 rows=22858 width=95)  
    (actual rows=23296 loops=1)
```

```
reltuples           | 16480  
most_common_vals    | {Berlin,...}  
most_common_freqs   | {0.0.0108333,...}
```

$$2.10993e+06 * 0.0108333 = 22857.5046690$$



Underestimate

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Berlin'
                        AND state_name = 'Berlin';
```

QUERY PLAN

```
-----
Seq Scan on zip_codes  (cost=0.00..62783.89 rows=263 width=95)
                        (actual rows=23296 loops=1)
  Filter: (((place_name)::text = 'Berlin'::text)
           AND ((state_name)::text = 'Berlin'::text))
  Rows Removed by Filter: 2086144
Planning Time: 0.157 ms
Execution Time: 220.953 ms
```




$$P(A \& B) = P(A) * P(B)$$



```
SELECT * FROM zip_codes
      WHERE place_name = 'Berlin'
      AND state_name = 'Berlin';
```

```
P(place_name = 'Berlin' & state_name = 'Berlin')
= P(city = 'Berlin') * P(state_name = 'Berlin')
= 0.0108333 * 0.0115
= 0.00012458295
```

$$2.10993e+06 * 0.00012458295 = 262.86$$



Underestimate

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Berlin'
        AND state_name = 'Berlin';
```

QUERY PLAN

```
-----
Seq Scan on zip_codes  (cost=0.00..62783.89 rows=263 width=95)
    (actual rows=23296 loops=1)
    Filter: (((place_name)::text = 'Berlin'::text)
        AND ((state_name)::text = 'Berlin'::text))
    Rows Removed by Filter: 2086144
Planning Time: 0.174 ms
Execution Time: 217.859 ms
```



Overestimate

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Berlin'
        AND state_name != 'Berlin';
```

QUERY PLAN

```
Seq Scan on zip_codes  (cost=0.00..62783.89 rows=22595 width=95)
    (actual rows=0 loops=1)
  Filter: (((state_name)::text <> 'Berlin'::text)
    AND ((place_name)::text = 'Berlin'::text))
 Rows Removed by Filter: 2109440
Planning Time: 0.143 ms
Execution Time: 232.766 ms
```



Correlated columns

- Attribute Value Independence Assumption (AVIA)
 - may result in wildly inaccurate estimates
 - both underestimates and overestimates
- consequences
 - poor scan choices (Seq Scan vs. Index Scan)
 - poor join choices (Nested Loop)



Poor scan choices

Index Scan using orders_city_idx on orders
(cost=0.28..185.10 **rows=90** width=36)
(actual **rows=12248237** loops=1)

Seq Scan using on orders
(cost=0.13..129385.10 **rows=12248237** width=36)
(actual **rows=90** loops=1)



Poor join choices

- > Nested Loop (... rows=90 ...) (... **rows=12248237** ...)
- > Index Scan using orders_city_idx on orders
(cost=0.28..185.10 **rows=90** width=36)
(actual **rows=12248237** loops=1)
...
- > Index Scan ... (... **loops=12248237**)



Poor join choices

```
-> Nested Loop (... rows=90 ...) (... rows=12248237 ...)
    -> Nested Loop (... rows=90 ...) (... rows=12248237 ...)
        -> Nested Loop (... rows=90 ...) (... rows=12248237 ...)
            -> Index Scan using orders_city_idx on orders
                (cost=0.28..185.10 rows=90 width=36)
                (actual rows=12248237 loops=1)
                ...
            -> Index Scan ... (... loops=12248237)
        -> Index Scan ... (... loops=12248237)
    -> Index Scan ... (... loops=12248237)
-> Index Scan ... (... loops=12248237)
```




functional dependencies (WHERE)



Functional Dependencies

- value in column A determines value in column B
- trivial example: primary key determines everything
 - zip code \rightarrow {place, community, county, state}
 - 89346 \rightarrow {Bibertal, Landkreis Günzburg, Swabia, Bayern}
- other dependencies:
 - place \rightarrow community
 - community \rightarrow county
 - county \rightarrow state



CREATE STATISTICS

```
CREATE STATISTICS s (dependencies)
  ON place_name, state_name, county_name FROM zip_codes;
ANALYZE zip_codes;
SELECT stxdependencies FROM pg_statistic_ext WHERE stxname = 's';
```

stxdependencies

```
-----
{"3 => 4": 0.918083, "3 => 6": 0.954369,
 "4 => 6": 0.689745, "6 => 4": 0.310255,
 "3, 4 => 6": 0.994842,
 "3, 6 => 4": 0.954672}
(1 row)
```



place → state: 0.918083 = d

$$P(\text{place} = \text{'Berlin'} \ \& \ \text{state} = \text{'Berlin'}) = \\ P(\text{place} = \text{'Berlin'}) * [d + (1-d) * P(\text{state} = \text{'Berlin'})]$$
$$16480 * 0.011 * (0.918 + (1-0.918) * 0.012) = 166.6$$



Underestimate : fixed

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Berlin'
        AND state_name = 'Berlin';
```

QUERY PLAN

```
-----
Seq Scan on zip_codes  (cost=0.00..62778.85 rows=22573 width=95)
    (actual rows=23296 loops=1)
   Filter: (((place_name)::text = 'Berlin'::text)
            AND ((state_name)::text = 'Berlin'::text))
  Rows Removed by Filter: 2086144
Planning Time: 0.228 ms
Execution Time: 216.878 ms
```



Overestimate #1: not fixed :-)

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Berlin'
        AND state_name != 'Berlin';
```

QUERY PLAN

```
-----
Seq Scan on zip_codes  (cost=0.00..62783.89 rows=22595 width=95)
    (actual rows=0 loops=1)
    Filter: (((state_name)::text <> 'Berlin'::text)
        AND ((place_name)::text = 'Berlin'::text))
    Rows Removed by Filter: 2109440
Planning Time: 0.152 ms
Execution Time: 240.476 ms
```

Functional dependencies only work with equalities.



Overestimate #2: not fixed :-)

```
EXPLAIN (ANALYZE, TIMING off)
SELECT * FROM zip_codes WHERE place_name = 'Berlin'
        AND state_name = 'Bayern';
```

QUERY PLAN

```
-----
Seq Scan on zip_codes  (cost=0.00..62778.85 rows=22768 width=95)
    (actual rows=0 loops=1)
   Filter: (((place_name)::text = 'Berlin'::text)
            AND ((state_name)::text = 'Bayern'::text))
  Rows Removed by Filter: 2109440
Planning Time: 0.074 ms
Execution Time: 217.444 ms
```

The queries need to respect the functional dependencies.



ndistinct (GROUP BY)



```
EXPLAIN (ANALYZE, TIMING off)
SELECT 1 FROM zip_codes GROUP BY community_name;
```

QUERY PLAN

```
HashAggregate  (cost=57509.08..57513.07 rows=399 width=23)
    (actual rows=400 loops=1)
    Group Key: community_name
    -> Seq Scan on zip_codes  (cost=0.00..52234.26 rows=2109926 width=19)
        (actual rows=2109440 loops=1)
```

Planning Time: 0.087 ms

Execution Time: 445.727 ms



```
SELECT attname, n_distinct
FROM pg_stats WHERE tablename = 'zip_codes';
```

attname	n_distinct
country_code	1
postal_code	7798
place_name	13326
state_name	16
stat_code	16
county_name	19
county_code	20
community_name	399
community_code	401
...	...

(12 rows)



```
EXPLAIN (ANALYZE, TIMING off)
```

```
SELECT 1 FROM zip_codes GROUP BY state_name, county_name, community_name;
```

QUERY PLAN

```
-----  
Group  (cost=418019.55..439114.20 rows=121296 width=55)
```

```
    (actual rows=400 loops=1)
```

```
    Group Key: state_name, county_name, community_name
```

```
    -> Sort  (cost=418019.55..423293.22 rows=2109465 width=51)
```

```
        (actual rows=2109440 loops=1)
```

```
        Sort Key: state_name, county_name, community_name
```

```
        Sort Method: external merge  Disk: 102160kB
```

```
        -> Seq Scan on zip_codes (cost=0.00..52229.65 rows=2109465 width=51)
```

```
            (actual rows=2109440 loops=1)
```

```
Planning Time: 0.276 ms
```

```
Execution Time: 3100.593 ms
```

```
(8 rows)
```



```
ndistinct(state, county, community)
=
ndistinct(state) * ndistinct(county) * ndistinct(communit
```

$$16 * 19 * 399 = 121296$$



```
CREATE STATISTICS s (ndistinct)
  ON state_name, county_name, community_name
  FROM zip_codes;
ANALYZE zip_codes;
```

```
SELECT stxndistinct FROM pg_statistic_ext;
```

```
      stxndistinct
```

```
-----
```

```
{ "4, 6": 31, "4, 8": 399,
  "6, 8": 399, "4, 6, 8": 397 }
(1 row)
```



```
EXPLAIN (ANALYZE, TIMING off)
```

```
SELECT 1 FROM zip_codes GROUP BY state_name, county_name, community_name;
```

QUERY PLAN

```
-----  
HashAggregate (cost=68051.18..68055.15 rows=397 width=55)
```

```
    (actual rows=400 loops=1)
```

```
    Group Key: community_name, state_name, county_name
```

```
    -> Seq Scan on zip_codes (cost=0.00..52229.96 rows=2109496 width=51)
```

```
        (actual rows=2109440 loops=1)
```

```
Planning Time: 0.243 ms
```

```
Execution Time: 656.465 ms
```

```
(5 rows)
```



ndistinct

- the “old behavior” was defensive
 - unreliable estimates with multiple columns
 - HashAggregate can’t spill to disk (OOM)
 - rather than crash do Sort+GroupAggregate (slow)
- ndistincts coefficients
 - make multi-column ndistinct estimates more reliable
 - reduced danger of OOM
 - large tables + GROUP BY multiple columns



Future Improvements

- additional types of statistics
 - MCV lists, histograms, ...
- statistics on expressions
 - currently only simple column references
 - alternative to functional indexes
- improving join estimates
 - using MCV lists
 - special multi-table statistics (syntax already supports it)



Questions?

Tomas Vondra

tomas.vondra@2ndquadrant.com
tomas@pgaddict.com



@fuzzycz



Questions?

Tomas Vondra

tomas.vondra@2ndquadrant.com

tomas@pgaddict.com



@fuzzycz